

SPARCLE · BOLT + AEIRA

The 2026 Agent-Orchestrator Choice

A side-by-side read of Bolt + Aeira against Google Gemini Spark, OpenAI Codex, Anthropic Claude Cowork, Microsoft Copilot Studio, and Glean Agents, after Google I/O 2026.

– THE QUESTION THIS DOCUMENT ANSWERS –

Who plays the orchestrator role across your SaaS stack: Google, OpenAI, Anthropic, Microsoft, or you?

Google I/O 2026 made it explicit. Gemini is positioned as the Client; your SaaS apps become data servers. OpenAI Codex, Anthropic Claude Cowork, and Microsoft Copilot Studio each assert the same role on different terms. Bolt + Aeira is the only product that keeps the orchestrator role inside your perimeter, on your LLM choice, on your audit trail.

WHAT THIS COVERS

15-row capability matrix across 7 platforms, anchored to public sources as of May 2026, plus a wedge analysis of what holds, what eroded, and what emerged post Google I/O 2026.

WHAT THIS DOES NOT

This is not a recommendation against the named vendors. Bolt + Aeira coexists with most of them. The matrix is to clarify where Sparcle is uniquely positioned, not to disparage incumbents.

HOW TO READ

Page 2 for the strategic frame. Pages 3 to 4 for the capability matrix. Pages 5 to 6 for the wedges. Page 7 for TCO at 1,000 seats. Page 8 for the pilot ask.

§02 The 2026 landscape map

Every flagship agent product on the market today sits in one of four quadrants. The axes are deployment posture (inside customer perimeter vs vendor-hosted) and LLM posture (single-vendor lock vs BYO LLM). Bolt + Aeira is alone in the top-right.

	Single-vendor LLM	BYO LLM (multi-vendor)
Inside customer perimeter	Empty quadrant No incumbent ships a perimeter deployment locked to a single LLM. The combination doesn't exist commercially in 2026.	Bolt + Aeira (Sparcle) Deploys via Helm, docker-compose, or sidecar inside customer VPC, on-prem, air-gap, or sovereign cloud. Customer chooses LLM per privacy tier. Customer pays LLM vendor directly. Zero Sparcle inference markup.
Vendor-hosted SaaS	Single-LLM SaaS agents Gemini Spark (Google · Gemini only), Codex (OpenAI · GPT only), Claude Cowork (Anthropic · Claude only), Microsoft 365 Copilot (Azure OpenAI bound). Each is a vertical play by the model vendor.	Cross-vendor SaaS aggregators Glean Agents routes across LLM choices but the data still flows through Glean's cloud. Copilot Studio supports limited model flexibility for built agents but the M365 surface is Azure OpenAI bound.

Why the empty quadrant matters

The empty top-left quadrant is structurally empty. Google, OpenAI, Anthropic, and Microsoft each have a financial reason to bind their agent product to their own LLM. None of them can ship a customer-perimeter deployment locked to a competitor's model. So the top-left will likely stay empty.

The top-right is also structurally hard. It requires an orchestrator that is model-agnostic by architecture, plus governance built for customer-controlled key material, plus connectors that are not tied to a single vendor's identity stack. Bolt + Aeira was designed for this quadrant from the first commit.

Strategic frame. The five named incumbents are competing for the same orchestrator role across your stack. Each wants their agent to be the Client and your other SaaS to be data servers. Bolt + Aeira is the product that lets that role stay with you.

What changed at Google I/O 2026 (May 19 to 20)

- **Gemini Spark** launched as a 24/7 personal agent running on dedicated GCE virtual machines per user. AI Ultra tier at USD 100 per user per month. Direct overlap with the launcher and watches surface in Bolt.
- **Antigravity 2.0** shipped as a macOS desktop agent harness with sub-agents, hooks, and asynchronous task primitives. Direct overlap with the bolt-native desktop launcher.
- **Search Agents** launched as 24/7 web monitors with Generative UI. Direct overlap with Aeira watches.
- **Google-managed Workspace MCP servers** shipped, with Gemini-as-Client framing. Either an opportunity (Bolt consumes them) or a moat (Workspace stays Gemini-orchestrated). Open question.
- **SynthID coalition** signed OpenAI, NVIDIA, Kakao, and ElevenLabs onto content credentials. A content-trust moat that Bolt should consider participating in for AI-generated artifacts.

§03 Capability matrix · part 1

Fifteen capability dimensions across seven platforms. Bolt + Aeira capabilities are anchored to shipped code paths. Competitor capabilities are anchored to vendor or press sources as of May 2026 (see page 8 footnotes).

Dimension	Bolt + Aeira	Gemini Spark / Antigravity	OpenAI Codex	Claude Cowork	MS Copilot Studio	Glean Agents
Deploys inside customer perimeter (on-prem · VPC · air-gap · sovereign)	✓ Helm · compose · sidecar · air-gap	✗ Google Cloud SaaS only	✗ OpenAI Cloud SaaS	✗ Anthropic Cloud SaaS	✗ Microsoft tenant	✗ Glean Cloud
Data never reaches vendor servers	✓ No Sparcle hosted path	✗	✗	✗	✗	✗
BYO LLM (any vendor, any model)	✓ OpenAI · Anthropic · Bedrock · NIM · Ollama · OpenAI-compatible	✗ Gemini only	✗ GPT only	✗ Claude only	✗ Azure OpenAI bound	~ Limited model choice
Zero inference markup (customer pays LLM vendor directly)	✓ Flat platform fee	✗ USD 100/seat covers compute	✗ Per-seat + token markup	✗ Per-seat + token markup	✗ USD 30/seat + Studio USD 200/mo tenant	✗ ~USD 50/user + AI add-on + support
PII masking at the LLM-call boundary (deterministic pseudonyms)	✓ HMAC-SHA256 · 12 default patterns + custom	✗ "Doesn't read indiscriminately" (no boundary contract)	✗ Tenant isolation only	✗ Tenant isolation only	✗ Tenant isolation only	✗ Tenant isolation only
Cryptographic audit chain (auditor-verifiable without vendor)	✓ Sealed Merkle · Ed25519 · standalone verifier binary	✗ Google's own logs	✗ OpenAI's own logs	✗ Anthropic's own logs	✗ M365 audit log scope only	✗ Glean's own logs
4-axis ACL (country · office · role-band · employment-type)	✓ Patent-pending · GIN-indexed	✗	✗	✗	✗ RBAC only	~ Permissions-aware knowledge graph
Per-region key residency + sovereignty enforcement	~ Per-region keyrings (Wave 3 pilot deliverable)	✗ Google-controlled regions	✗ OpenAI-controlled	✗ Anthropic-controlled	~ Microsoft sovereign clouds available	✗ Glean-controlled

✓ =shipped or vendor-confirmed available · ~ =partial · pilot-deliverable · or limited surface · ✗ =not available in product as of May 2026. Capability matrix continues on next page.

§04 Capability matrix · part 2

Continuing rows 9 through 15. The bottom seven dimensions are where the 2026 incumbents have moved fastest. Bolt + Aeira holds the orchestration depth; competitors have caught up on table-stakes (MCP, A2A, desktop, agents).

Dimension	Bolt + Aeira	Gemini Spark / Antigravity	OpenAI Codex	Claude Cowork	MS Copilot Studio	Glean Agents
Cross-OS desktop-native (macOS · Windows · Linux)	✓ mac · win · linux + PWA + extension	~ Antigravity macOS; Windows pending	✓ Codex.app on macOS + Windows	~ Cloud + web; Claude Code CLI	~ Windows-first	✗ Browser only
SaaS connector count (in-house / total reachable)	✓ 103 in-house + ~350 via Airbyte	~ 8 Workspace apps + Google MCP servers	~ 90+ plugins / MCP bundles	~ Gmail · DocuSign · Clay · FactSet · S&P · MSCI · PitchBook	~ Microsoft-first; Graph Connectors for the rest	✓ 100+ data connectors
Local OS execution (terminal · files · clipboard · apps) with sandbox	✓ 29+ tools · 4 sandbox tiers	~ Antigravity harness (dev workflows)	✓ Computer-use (sees · clicks · types)	✗ Claude Code is CLI scope only	✗ No general local execution	✗
MCP host (consume) + MCP server (expose itself)	✓ Bidirectional + 10+ auth strategies	✓ Google-managed Workspace MCP servers	✓ MCP via plugin marketplace	✓ MCP host + Managed Agents	~ Studio agents have MCP	✓ MCP host (consume only)
A2A cross-vendor agent delegation	✓ Full A2A · agent.json · durable tasks	✓ A2A protocol steward · 150+ orgs	~ Not surfaced as A2A endpoint	~ A2A capable (Anthropic adopted)	~ Work IQ A2A in public preview (Apr 2026)	~ Sub-agent orchestration internal only
Native automation triggers (cron · webhook · file-watch · etc.)	✓ 12 trigger types built-in	✓ Spark Schedules	~ Some scheduling	~ Scheduled agents	~ Power Automate (separate license)	~ Watch-style alerts
Per-role LLM gating + per-team MCP allowlist + audit access control	✓ 50+ admin endpoints · 18 RBAC permissions	~ Workspace admin only	~ Admin marketplaces	~ Managed Agents governance	✓ Enterprise-grade Studio admin	✓ Glean admin console

Read the matrix this way. The first eight rows on page 3 are deployment, data-flow, and governance dimensions where Bolt + Aeira holds structurally. The seven rows on this page are surface dimensions where the incumbents have caught up. The structural rows are durable; the surface rows are not what the comparison rests on.

Three quick observations

- **MCP and A2A are table-stakes now.** Six of seven products have MCP host support; four have some A2A surface. The 2025 differentiator has commoditized.
- **Desktop is contested.** Codex.app and Antigravity ship desktop apps; Bolt's cross-OS lead is real but not unique on the macOS axis.
- **Connectors are converging.** Anthropic added FactSet, S&P Capital IQ, MSCI, Pitchbook between February and May 2026. The connector-count game is now about quality of permission-aware coverage, not raw count.

§05 Three wedges that hold

After Google I/O 2026, three Bolt + Aeira capability vectors remain structurally unique. Each is grounded in customer-visible architecture, not marketing.

Wedge 1 · Deployment inside customer perimeter

Helm, docker-compose, or sidecar deployments run inside the customer's VPC, on-prem data center, air-gapped environment, or sovereign cloud (AWS, Azure, GCP regional / sovereign tiers). No data flows to Sparcle. The audit chain is verifiable by the customer's own auditor using a standalone binary.

None of the five named incumbents ships a customer-controlled deployment for their flagship agent. Google Gemini Spark runs on dedicated GCE virtual machines owned by Google. OpenAI Codex runs in OpenAI Cloud. Claude Cowork runs in Anthropic Cloud. Microsoft 365 Copilot runs in the Microsoft tenant. Glean Agents runs in Glean Cloud. The financial model of each ties the agent to that vendor's infrastructure; this is structural.

Wedge 2 · BYO LLM with privacy-tier routing

Bolt's LLM layer accepts OpenAI, Anthropic, Google, Bedrock, NVIDIA NIM, Ollama, and any OpenAI-compatible endpoint. Customer policy decides which model handles which class of data, by privacy tier: frontier, regional, open-weight, on-device. Customer pays the LLM vendor directly; Sparcle takes zero inference markup.

Each of the five incumbents has a financial incentive to bind the agent to its own LLM. Spark is Gemini-only. Codex is GPT-only. Cowork is Claude-only. Copilot Studio is Azure OpenAI bound. Glean offers limited model choice but takes margin on the AI add-on. Switching is contractual and operational, not configuration.

Wedge 3 · Aeira cryptographic governance substrate

The Aeira data plane provides three governance primitives no incumbent matches today.

- **4-axis ACL pre-filter** on country, office, role-band, and employment-type at index time. Patent-pending, three provisional filings.
- **Sealed Merkle audit chain**, Ed25519-signed, KMS-mediated. Customer-side auditor verifies a tamper-evident chain using a standalone binary without Sparcle's cooperation.
- **Crypto-shred** via per-tenant KEK deletion. GDPR Article 17 and HIPAA right-to-erasure satisfied in one primitive.

Glean's permissions-aware knowledge graph is the closest analog and is the most credible competitor on the governance axis. Glean still routes through Glean Cloud; the audit boundary is theirs, not yours.

One sentence to internalize. Bolt + Aeira is the only product on this comparison where the customer owns the deployment, the LLM choice, and the audit trail simultaneously. Removing any one of the three would put us back in a quadrant the incumbents already occupy.

§06 Two wedges that emerged from I/O 2026

Google's I/O 2026 framing made two further wedges visible. Both were latent in Bolt's architecture before the keynote; the keynote made them sellable.

Wedge 4 · Orchestrator role inside the customer's perimeter

At I/O 2026, Google made the strategic frame explicit. Gemini is the Client; SaaS apps including third-party MCP servers become data servers. The same framing fits OpenAI Codex (Codex is the Client), Anthropic Claude Cowork (Cowork is the Client), and Microsoft Copilot Studio (Copilot is the Client). The pattern across all four incumbents is identical: claim the orchestrator role.

Bolt + Aeira's counter-framing is that the orchestrator role belongs inside the customer's perimeter, with the customer's identity, policy, and audit boundary owning it. Bolt is the Client; the customer's MCP servers (including Google's Workspace MCP, if openly callable) are data servers; the customer's LLM choice is policy, not vendor lock.

Use this in prospect conversations. "Do you want Google, OpenAI, Anthropic, or Microsoft to be the orchestrator across your stack, or do you want that role to stay inside your perimeter, with your identity, your policy, and your audit trail?"

Wedge 5 · Policy-driven LLM routing per privacy tier

Each I/O 2026 incumbent picks the model for you. Spark picks Gemini 3.5 Flash. Codex picks GPT-5.5. Cowork picks Claude. Copilot picks Azure OpenAI.

Bolt's LLM router lets the customer's policy decide model selection per request, based on the data's privacy tier. A query that touches unrestricted public data may route to a frontier model; a query that touches PII may route to a regional or on-device model. The routing decision is auditable and reversible. Sparcle has no economic reason to push any one model.

What this means for the pitch

Stop benchmarking Bolt against per-seat pricing of any one incumbent. The product is not a per-seat replacement for Spark or Copilot. It is the infrastructure that keeps the orchestrator role and the model-routing decision inside the customer's boundary, while the customer continues using whichever incumbents fit specific workflows (Copilot inside M365, Claude Cowork for finance templates, Spark for personal scheduling).

What is honestly not a Bolt + Aeira wedge

- **Consumer 24/7 personal-agent surface.** Spark wins. Bolt is not a consumer product.
- **Vertical agent template library.** Claude Cowork ships 10 finance, 12 legal, and growing HR templates. We have personas (security, finance shipped) but the vertical-template depth is theirs.
- **Developer IDE harness.** Antigravity and Codex.app are dev-tooling targets. Bolt is an enterprise orchestrator, not an IDE.
- **Web UX polish.** Glean has years of polish. We will be three to six months behind through 2026.

§07 TCO comparison at 1,000 seats

Three-year total cost of ownership for a 1,000-seat deployment, using public list pricing where available, May 2026. Numbers exclude implementation, training, and customer infrastructure. Bolt + Aeira numbers use the Bolt Bundled tier (USD 60 per seat per month, customer self-hosted, BYO LLM with BYOK).

Product	Y1 platform	Y1 LLM	Y1 total	3yr total
Bolt + Aeira (Bundled, BYO LLM)	\$720,000	\$300,000 (customer-direct)	\$1,020,000	\$3,060,000
Gemini Spark (AI Ultra)	\$1,200,000 (USD 100/seat)	included	\$1,200,000	\$3,600,000
OpenAI Codex (per-seat est.)	~\$720,000 (USD 60/seat est.)	+ token usage (est. \$250K)	~\$970,000	~\$2,910,000
Claude Cowork (per-seat est.)	~\$720,000 (USD 60/seat est.)	+ token usage (est. \$300K)	~\$1,020,000	~\$3,060,000
MS Copilot + Studio	\$360,000 (USD 30/seat) + \$2,400 tenant	included	\$362,400	\$1,087,200
Glean Agents (with AI add-on)	\$600,000 platform + \$180,000 AI + 10% support = \$858,000	included	\$858,000	\$2,574,000

Codex and Cowork per-seat numbers are estimates; both vendors negotiate enterprise pricing privately. Token-usage estimates assume moderate-intensity enterprise agent use (3 to 5 million tokens per seat per month). LLM cost line for Bolt + Aeira assumes customer-direct OpenAI / Anthropic / Bedrock at observed-blended rates.

What the TCO does and does not say

Bolt + Aeira is not the cheapest line by raw three-year TCO. Microsoft Copilot Studio is cheaper because Microsoft subsidizes Copilot inside the M365 bundle, and Glean's AI tier sits below Spark's all-in price.

What Bolt + Aeira's number actually buys is the structural posture: customer-perimeter deployment, BYO LLM with zero markup, auditor-verifiable Merkle chain, 4-axis ACL, crypto-shred. For regulated industries where deployment topology is the compliance conversation, the Microsoft and Glean rows are not adoptable at all. For non-regulated buyers, Microsoft and Glean are viable; Bolt + Aeira competes on governance depth, not raw cost.

Cost lines competitors carry that Bolt + Aeira does not

- **Per-user VM cost.** Spark runs on dedicated GCE virtual machines per user. The USD 100 per seat per month price floor is structural, not promotional.
- **Inference markup.** Codex, Cowork, Copilot, and Spark all bake LLM cost into the per-seat price with vendor margin. Bolt's customer pays the LLM vendor directly with no Sparcle markup.
- **Mandatory add-on bundles.** Glean's mandatory 10 percent support fee, Copilot Studio's USD 200 tenant fee, AI Ultra's all-in subscription model. Bolt's pricing is flat.

Where Bolt + Aeira wins on cost. Customers running 10K to 100K seats, where the per-user VM model of Spark or the per-seat plus token markup of Codex / Cowork compound. At scale, Bolt's flat platform fee plus customer-direct LLM cost is the model that does not scale linearly with seat count.

§08 Decision frame + sources

A short decision frame for the buyer, then the public sources every claim in this document is anchored to.

Decision frame for the buyer

Choose Bolt + Aeira when

- Deployment topology is the compliance conversation (regulated, sovereign, air-gap)
- You want the orchestrator role to stay inside your perimeter
- You want LLM choice as policy, not vendor lock
- You want auditor-verifiable governance without depending on a vendor's word
- You operate across multiple SaaS vendors and need cross-source orchestration

Choose an incumbent when

- You operate inside a single vendor's stack (pure M365 or pure Workspace shop)
- You have no regulated data and no sovereignty pressure
- You want a consumer 24/7 personal agent (Spark is the right product)
- You want vertical-template depth in finance or legal (Cowork is the right product)
- You are buying for a developer team only (Codex or Antigravity is the right product)

The 90-day pilot ask

Bolt + Aeira pilots run 90 days inside the customer's infrastructure. Day 1 to 14 is deployment and ingest. Day 15 to 60 is workflow validation against the customer's actual data. Day 61 to 90 produces the case-study artifact and the renewal decision. Pilot is no-obligation with a 30-day tear-down clean-exit clause. Founding-customer pricing locks for 24 months on conversion to annual.

Sources (public, as of May 2026)

1. **Gemini Spark · Antigravity 2.0 · Search Agents.** gemini.google/overview/agent/spark/ · blog.google/innovation-and-ai/technology/developers-tools/google-io-2026-collection/ · cloud.google.com/blog/topics/developers-practitioners/io26-news-for-agent-developers-on-google-cloud · 9to5google.com/2026/05/19/google-io-2026-news/
2. **OpenAI Codex (app + plugin marketplace + computer-use).** openai.com/index/introducing-the-codex-app/ · openai.com/codex/ · developers.openai.com/codex/changelog · cloudwars.com/ai/openai-moves-closer-to-super-app-vision-as-codex-expands
3. **Anthropic Claude Cowork + Managed Agents + Finance Templates.** anthropic.com/news/finance-agents · pasqualepillitteri.it/en/news/755/anthropic-managed-agents-cowork-ga-april-9-2026 · techcrunch.com/2026/02/24/anthropic-launches-new-push-for-enterprise-agents
4. **Microsoft 365 Copilot Studio + Work IQ A2A.** microsoft.com/en-us/microsoft-copilot/blog/copilot-studio/new-and-improved-agent-governance-intelligent-workflows-and-connected-app-experiences/ · learn.microsoft.com/en-us/power-platform/release-plan/2026wave1/microsoft-copilot-studio/
5. **Glean Agents.** glean.com/glean-vs-copilot · futurumgroup.com/insights/glean-doubles-arr-to-200m-can-its-knowledge-graph-beat-copilot
6. **A2A protocol.** developers.googleblog.com/en/a2a-a-new-era-of-agent-interopability/ · github.com/a2aproject/A2A · stellagent.ai/insights/a2a-protocol-google-agent-to-agent
7. **Bolt + Aeira claims.** Anchored to shipped code paths in the Sparcle codebases (bolt-api, bolt-pwa, bolt-native, aeira-rs, bolt-common, bolt-license) and to the canonical claim-hygiene file at sparcle-LLC/business/outreach/playbook/02-claim-hygiene.md.

Sovereign enterprise AI, inside your perimeter, with your LLM, on your audit trail.

sparcle.app/strategy · bolt@sparcle.app · sparcle-agent-landscape-2026.pdf · v1 · 2026-05 · Confidential