

Value & Cost — the honest math.

Sparcle Bolt Absolute (sovereign tier) + Aeira Enhanced · regulated 1,000-seat enterprise · self-hosted · 1-year illustrative scenario.

Annual cost — fully disclosed

Sparcle license + customer-side infrastructure

NET-NEW SPARCLE LICENSE	
Bolt Absolute · per seat / month	\$360,000
\$30/seat × 1,000 × 12 mo. Absolute = self-hosted, BYO LLM, full data sovereignty — the regulated-fit tier.	
Aeira Enhanced · fixed monthly	\$59,988
\$4,999/mo × 12. Production-scale: 100K records, Private Instance License, SLA guarantees. Fixed cost, not per-seat.	
Founding-customer discount · 40%, 24-mo lock	−\$167,995
First 50 customers. Applies to Bolt + Aeira list price.	
Net-new annual Sparcle spend	\$251,993
CUSTOMER-SIDE INFRASTRUCTURE · REAL MONEY YOU PAY	
Compute · Bolt + Aeira VPC nodes	~\$60,000
~6 mid-size VMs: API · ingest · scheduler · Postgres HA · NATS · Redis.	
Storage · Postgres + S3 + audit WORM	~\$20,000
~500GB Postgres + ~2TB S3 + Object Lock for audit chain.	
LLM tokens · your existing BYO LLM contract	~\$80,000
~6B tokens/yr at blended \$5-15/M. GPT-4o, Claude 3.5, or self-hosted Llama.	
Customer-side annual operating cost	~\$160,000
These costs are real money but typically fit inside your existing AWS / Azure / GCP commit + LLM enterprise contract — <1% of existing infra budget for a multi-million-dollar spender. No new vendor relationships.	
Total annual customer outlay (Sparcle + infra)	~\$411,993

What this calculation is — and is not

This is the framework, not a guarantee. The productivity number is industry-benchmarked from competitor case studies and academic studies, not Sparcle-measured customer outcomes. The breach + PII numbers are expected-value calculations using public-report inputs and conservative midpoints. **Your scenario will differ.** Ask us for a tenant-specific calculator — plug in your seats, your LLM rates, your data-sensitivity profile, and fully-loaded labor rate. Categorical differentiators add architectural value that doesn't fit on a dollar line; cost is competitive at Absolute tier (~40% below Glean per-seat).

Annual value created or avoided

Conservative midpoints · sourced from public reports

VALUE CREATED + AVOIDED	
Productivity acceleration	\$10,000,000
Industry benchmark: 80-150 hrs saved per knowledge worker per year from cross-source enterprise AI search (Glean published cases; McKinsey AI 2024 productivity studies). Conservative midpoint: 100 hrs × \$100/hr blended labor × 1,000 seats.	
Data-breach avoidance (expected value)	\$244,000
\$4.88M avg breach cost (IBM 2024) × 5% expected annual GenAI-incident probability (industry range 3-10% for unprotected adopters; conservative midpoint). Sparcle's in-perimeter posture eliminates vendor-side breach surface.	
PII / NPI / PHI leak avoidance	\$500,000
Expected-value range \$250K-\$2M/yr depending on data-sensitivity profile. ~5,000 PII tokens/worker/yr to LLMs (Cyberhaven 2024). Conservative midpoint shown.	
Audit-readiness cost avoidance	\$120,000
Standalone cryptographic verifier reduces SOX / SEC / HHS evidence-gathering effort. Conservative: ~800 staff hours × \$150/hr.	
Third-party-vendor-risk-review avoidance	\$50,000
Sparcle deploys in your perimeter — no new sub-processor. ~\$50K of CISO + procurement effort per year saved.	
Total annual value (conservative midpoints)	\$10,914,000

Scenario assumptions

1,000 knowledge-worker seats · regulated buyer
Self-hosted in customer VPC or on-prem
Bolt **Absolute** tier (\$30/seat/mo) · Aeira Enhanced (\$4,999/mo)
BYO LLM via customer's existing contract
Founding-customer 40% discount applied (24-mo lock)

NET ANNUAL POSITION · ILLUSTRATIVE

Value created	\$10.91M
Customer total outlay	\$0.41M

~26x

value/outlay ratio in this scenario.
Range across conservative assumptions: ~20x to ~40x.

Plus — categorical differentiators not priced in this math

- 1. Inside your perimeter.** Data never reaches Sparcle. No SaaS competitor (Glean / Hebbia / Harvey / Copilot) can match this — at any price.
- 2. BYO LLM, zero markup.** Use the LLM you already contracted. Competitors bundle inference at 30-60% markup.
- 3. Standalone audit verifier.** Your auditor verifies our chain without our cooperation. Competitors ship vendor-attested logs.
- 4. No new sub-processors.** Your VRM register doesn't grow. Glean/Hebbia/Harvey each add a sub-processor under GDPR Art. 28.
- 5. Patent-pending 4-axis ACL.** country · office · role-band · employment-type. Nobody else indexes on these dimensions.
- 6. Source-code escrow.** Survives our shutdown. SaaS competitors' shutdown leaves you stranded mid-contract.

Sources + caveats

- \$4.88M avg breach cost — IBM Cost of a Data Breach 2024
- \$9.77M healthcare breach (14th yr #1) — IBM 2024
- 80-150 hrs/seat/yr productivity range — Glean published cases (Confluent, Pinterest); McKinsey State of AI 2024
- ~5,000 PII tokens/worker/yr — Cyberhaven 2024 GenAI usage study
- 3-10% expected GenAI breach probability — industry-aggregate range
- Sparcle pricing — sparcle.app/pricing (Bolt Absolute / Aeira Enhanced)
- Customer infra estimates conservative midpoints; vary by topology + LLM choice